

In this notebook I am going to try various pre-trained CNN models such as VGG16, InceptionV3, EfficientNet and ResNet along with a simple CNN baseline model using transfer learning method.

```
#importing libraries
import pandas as pd
import numpy as np
```

```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import warnings
warnings.filterwarnings("ignore")
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
df = pd.read_csv('/content/drive/MyDrive/Applied_ai/df_img.csv')
```

```
df.head()
```

	url	dank_or_not
0	https://i.redd.it/s044gyh084061.jpg	1
1	https://i.imgur.com/G2bnxQa.jpg	1
2	https://i.redd.it/08bjn1pan3061.jpg	1
3	https://i.redd.it/e8nptzrqs2061.jpg	1
4	https://i.redd.it/8dwnjgd962061.jpg	1

```
#saving only name of each images that was downloaded
df['url'] = df['url'].str.split("/").str[-1].str.replace('.jpg', '.png')
```

```
#splitting the data
from sklearn.model_selection import train_test_split
train_df, val_df = train_test_split(df, test_size=0.2, random_state=42, shuffle=True)
```

```
#creating train and test image generators and performing data augmentation on only
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
train_datagen = ImageDataGenerator(rescale=1./255,
                                   shear_range=0.2,
```

```

        zoom_range=0.2,
        rotation_range=45,
        horizontal_flip=True,
        vertical_flip=True)

test_datagen = ImageDataGenerator(rescale=1./255)

train_ImageGenerator = train_datagen.flow_from_dataframe(train_df, x_col = 'url',
test_ImageGenerator = test_datagen.flow_from_dataframe(val_df,x_col = 'url', y_col

Found 3996 validated image filenames.
Found 999 validated image filenames.

```

```

#importing libraries
import tensorflow as tf
from tensorflow.keras.utils import plot_model
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense, Flatten
from tensorflow.keras.layers import Conv2D, MaxPool2D, GlobalAveragePooling2D
from tensorflow.keras.metrics import Accuracy
accuracy = Accuracy()
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

#pre-steps to load tensorboard in colab
! wget https://bin.equinox.io/c/4VmDzA7iaHb/ngrok-stable-linux-amd64.zip > /dev/nu
! unzip ngrok-stable-linux-amd64.zip > /dev/null 2>&1

LOG_DIR = './log'
get_ipython().system_raw(
    'tensorboard --logdir {} --host 0.0.0.0 --port 6006 &'
    .format(LOG_DIR)
)

get_ipython().system_raw('./ngrok http 6006 &')

! curl -s http://localhost:4040/api/tunnels | python3 -c \
    "import sys, json; print(json.load(sys.stdin)['tunnels'][0]['public_url'])"

https://e89e956fa5f9.ngrok.io

```

▼ Baseline CNN model

This simple model is our baseline model which will be trained using image data present.

```

#Creating baseline model with two convolution layers
input = Input(shape=(224,224,3))

conv1 = Conv2D(64, kernel_size=3, activation='relu')(input)
pool1 = MaxPool2D(pool_size=(2,2))(conv1)

```

```
pool1 = MaxPool2D(pool_size=(2,2))(conv1)

conv2 = Conv2D(128, kernel_size=3, activation='relu')(pool1)
pool2 = MaxPool2D(pool_size=(2,2))(conv2)

flat = Flatten()(pool2)
dense1 = Dense(32, activation='relu')(flat)
output = Dense(1, activation='sigmoid')(dense1)

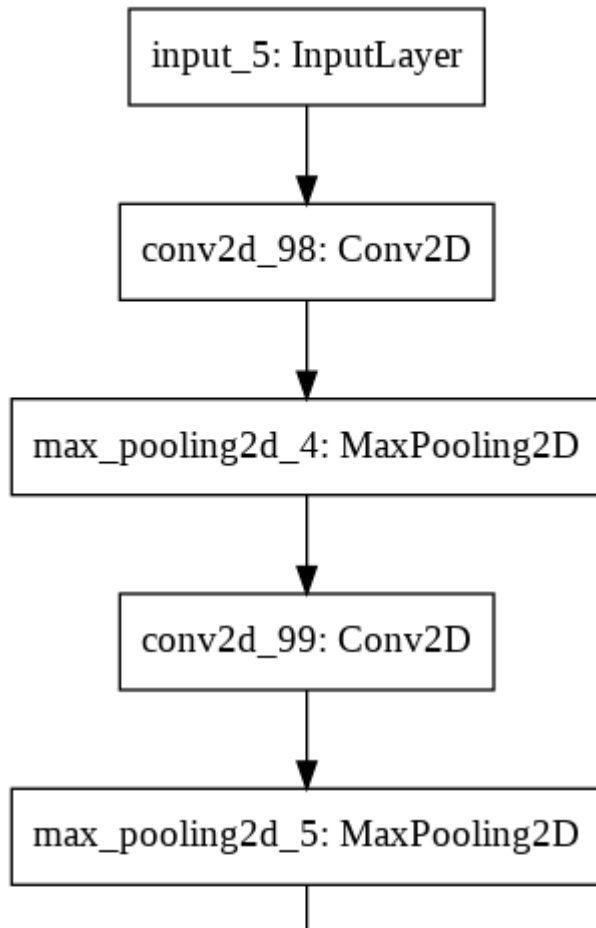
model = Model(inputs=input, outputs=output)

print(model.summary())
```

Model: "model_4"

Layer (type)	Output Shape	Param #
input_5 (InputLayer)	[(None, 224, 224, 3)]	0
conv2d_98 (Conv2D)	(None, 222, 222, 64)	1792
max_pooling2d_4 (MaxPooling2D)	(None, 111, 111, 64)	0
conv2d_99 (Conv2D)	(None, 109, 109, 128)	73856
max_pooling2d_5 (MaxPooling2D)	(None, 54, 54, 128)	0
flatten_2 (Flatten)	(None, 373248)	0
dense_7 (Dense)	(None, 32)	11943968
dense_8 (Dense)	(None, 1)	33
Total params: 12,019,649		
Trainable params: 12,019,649		
Non-trainable params: 0		
None		

```
plot_model(model)
```



```
%load_ext tensorboard
# Clear any logs from previous runs
!rm -rf ./log/
```

The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

```
| . _ _ |
```

```
#tensorboard callback to draw different metrics during training
tbCallBack = tf.keras.callbacks.TensorBoard(log_dir='./log', histogram_freq=0,
                                             write_graph=True,
                                             write_grads=True,
                                             write_images=True)
callback_list = [tbCallBack]

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `1

#compiling the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

#training the model
model.fit_generator(train_ImageGenerator, epochs = 20, validation_data=(test_Image
```

```
Epoch 1/20
125/125 [=====] - 129s 956ms/step - loss: 0.9685 - a
Epoch 2/20
125/125 [=====] - 126s 961ms/step - loss: 0.6931 - a
Epoch 3/20
```

```
125/125 [=====] - 141s 1s/step - loss: 0.6929 - accu
Epoch 4/20
125/125 [=====] - 126s 941ms/step - loss: 0.6927 - a
Epoch 5/20
125/125 [=====] - 129s 960ms/step - loss: 0.6923 - a
Epoch 6/20
125/125 [=====] - 126s 953ms/step - loss: 0.6935 - a
Epoch 7/20
125/125 [=====] - 126s 943ms/step - loss: 0.6920 - a
Epoch 8/20
125/125 [=====] - 125s 943ms/step - loss: 0.6924 - a
Epoch 9/20
125/125 [=====] - 127s 953ms/step - loss: 0.6910 - a
Epoch 10/20
125/125 [=====] - 126s 953ms/step - loss: 0.6913 - a
Epoch 11/20
125/125 [=====] - 126s 944ms/step - loss: 0.6943 - a
Epoch 12/20
125/125 [=====] - 127s 948ms/step - loss: 0.6909 - a
Epoch 13/20
125/125 [=====] - 126s 933ms/step - loss: 0.6902 - a
Epoch 14/20
125/125 [=====] - 125s 953ms/step - loss: 0.6906 - a
Epoch 15/20
125/125 [=====] - 140s 1s/step - loss: 0.6889 - accu
Epoch 16/20
125/125 [=====] - 126s 947ms/step - loss: 0.6914 - a
Epoch 17/20
125/125 [=====] - 125s 938ms/step - loss: 0.6909 - a
Epoch 18/20
125/125 [=====] - 126s 942ms/step - loss: 0.6902 - a
Epoch 19/20
125/125 [=====] - 127s 967ms/step - loss: 0.6893 - a
Epoch 20/20
125/125 [=====] - 126s 941ms/step - loss: 0.6899 - a
<tensorflow.python.keras.callbacks.History at 0x7fe243dcc390>
```



```
#drawing tensorboard
%tensorboard --logdir log
```

TensorBoard

SCALARS

GRAPHS

INACTIVE

- ☐ Show data download links
- ☐ Ignore outliers in chart scaling

Tooltip sorting method: **default** ▼

Smoothing



0.6


Horizontal Axis

STEP

RELATIVE

WALL


Runs



epoch_loss



epoch_loss
tag: epoch_loss



#predicting use trained baseline model and it will predict the probabilities for t
y_pred_b = model.predict(test_ImageGenerator)

validation

#converting probabilities to labels

y_pred_b_df = pd.DataFrame(y_pred_b).round().astype('int')

log

#plotting confusion matrix

cm_b = confusion_matrix(val_df['dank_or_not'], y_pred_b_df[0])

sns.heatmap(cm_b, annot=True, fmt='g')

plt.xlabel('Actual')

plt.ylabel('Predicted')

plt.show()

```
#getting true negative, false positive, false negative and true positive values fr
tn_b, fp_b, fn_b, tp_b = cm_b.ravel()
```

```
#computing sensetivity and specificity
sensitivity_b = (tp_b/(tp_b+fn_b)).round(4)
specificity_b = (tn_b/(tn_b+fp_b)).round(4)
print('sensitivity      : ',sensitivity_b)
print('specificity      : ',specificity_b)
print('accuracy         : ',accuracy(val_df['dank_or_not'], y_pred_b_df[0]))

sensitivity      :  0.508
specificity      :  0.511
accuracy         :  tf.Tensor(0.5095095, shape=(), dtype=float32)
```

```
model.save('baseline.h5')
```

▼ Applying VGG16

```
#importing VGG16 model
from tensorflow.keras.applications.vgg16 import VGG16
```

```
%load_ext tensorboard
# Clear any logs from previous runs
!rm -rf ./log2/
```

```
The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard
```

```
#defining VGG16 model without including top fully connected layers, to perform tra
vgg16_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224,
```

```
#setting layers of VGG16 as non-trainable
for layer in vgg16_model.layers:
    layer.trainable = False
```

```
#checking layers of VGG16
for i, layer in enumerate(vgg16_model.layers):
    print(i, layer.name, layer.trainable)
```

```
0 input_4 False
1 block1_conv1 False
2 block1_conv2 False
3 block1_pool False
4 block2_conv1 False
5 block2_conv2 False
6 block2_pool False
7 block3_conv1 False
8 block3_conv2 False
```

```

9 block3_conv3 False
10 block3_pool False
11 block4_conv1 False
12 block4_conv2 False
13 block4_conv3 False
14 block4_pool False
15 block5_conv1 False
16 block5_conv2 False
17 block5_conv3 False
18 block5_pool False

```

```
vgg16_model.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
Total params: 14,714,688		
Trainable params: 0		
Non-trainable params: 14,714,688		


```
#adding layers after non-trainable layers of VGG16
x = vgg16_model.output
x = (Conv2D(filters = 128, kernel_size = 7, strides=1, activation='relu', kernel_i
x = (Conv2D(filters = 128, kernel_size = 1, strides=1, activation='relu', kernel_i
x = (Flatten()))(x)
x = (Dense(1, activation='sigmoid'))(x)
```


```
model2 = Model(inputs=vgg16_model.input, outputs=x)
```

```
#tensorboard callback
```

```
tbCallBack2 = tf.keras.callbacks.TensorBoard(log_dir='./log2', histogram_freq=0,
                                              write_graph=True,
                                              write_grads=True,
                                              write_images=True)
```

```
callback_list2 = [tbCallBack2]
```

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `



```
#compiling the model
```

```
model2.compile(optimizer='adam', loss="binary_crossentropy", metrics=["accuracy"])
```

```
#training the model, only the layers below non-trainable layers will be trained
```

```
model2.fit_generator(train_ImageGenerator, epochs = 20, validation_data=(test_Imag
```

```
Epoch 1/20
125/125 [=====] - 129s 956ms/step - loss: 0.7722 - a
Epoch 2/20
125/125 [=====] - 128s 960ms/step - loss: 0.6982 - a
Epoch 3/20
125/125 [=====] - 133s 960ms/step - loss: 0.6740 - a
Epoch 4/20
125/125 [=====] - 129s 969ms/step - loss: 0.6823 - a
Epoch 5/20
125/125 [=====] - 129s 977ms/step - loss: 0.6767 - a
Epoch 6/20
125/125 [=====] - 132s 949ms/step - loss: 0.6743 - a
Epoch 7/20
125/125 [=====] - 145s 1s/step - loss: 0.6647 - accu
Epoch 8/20
125/125 [=====] - 127s 957ms/step - loss: 0.6753 - a
Epoch 9/20
125/125 [=====] - 127s 960ms/step - loss: 0.6644 - a
Epoch 10/20
125/125 [=====] - 128s 951ms/step - loss: 0.6743 - a
Epoch 11/20
125/125 [=====] - 128s 967ms/step - loss: 0.6706 - a
Epoch 12/20
125/125 [=====] - 133s 945ms/step - loss: 0.6710 - a
Epoch 13/20
125/125 [=====] - 127s 951ms/step - loss: 0.6605 - a
Epoch 14/20
125/125 [=====] - 142s 1s/step - loss: 0.6537 - accu
Epoch 15/20
125/125 [=====] - 128s 958ms/step - loss: 0.6570 - a
Epoch 16/20
```

```
125/125 [=====] - 130s 978ms/step - loss: 0.6533 - a
Epoch 17/20
125/125 [=====] - 142s 1s/step - loss: 0.6510 - accu
Epoch 18/20
125/125 [=====] - 142s 1s/step - loss: 0.6532 - accu
Epoch 19/20
125/125 [=====] - 127s 960ms/step - loss: 0.6567 - a
Epoch 20/20
125/125 [=====] - 148s 1s/step - loss: 0.6413 - accu
<tensorflow.python.keras.callbacks.History at 0x7fe244438590>
```



```
#generating tensorboard
%tensorboard --logdir log2
```

TensorBoard

SCALARS

GRAPHS


INACTIVE

```
#predicting values for test dataset
y_pred_v = model2.predict(test_ImageGenerator)
```

```
y_pred_v_df = pd.DataFrame(y_pred_v).round().astype('int')
```

Tooltip sorting

```
#plotting confusion matrix
cm_v = confusion_matrix(val_df['dank_or_not'], y_pred_v_df[0])
sns.heatmap(cm_v, annot=True, fmt='g')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



TOGGLE ALL RINGS

```
#getting true negative, false positive, false negative and true positive values fr
tn_v, fp_v, fn_v, tp_v = cm_v.ravel()
```

```
#computing sensetivity and specificity
sensitivity_v = (tp_v/(tp_v+fn_v)).round(4)
specificity_v = (tn_v/(tn_v+fp_v)).round(4)
print('sensitivity      : ',sensitivity_v)
print('specificity      : ',specificity_v)
print('accuracy         : ',accuracy(val_df['dank_or_not'], y_pred_v_df[0]))
```

```
sensitivity      :  0.4719
specificity      :  0.5269
accuracy         :  tf.Tensor(0.50784117, shape=(), dtype=float32)
```

```
#saving the model
model2.save('vgg16_model.h5')
```

▼ Applying Inception v3 model


```

281 batch_normalization_89 False
282 activation_89 False
283 conv2d_86 False
284 conv2d_90 False
285 batch_normalization_86 False
286 batch_normalization_90 False
287 activation_86 False
288 activation_90 False
289 conv2d_87 False
290 conv2d_88 False
291 conv2d_91 False
292 conv2d_92 False
293 average_pooling2d_8 False
294 conv2d_85 False
295 batch_normalization_87 False
296 batch_normalization_88 False
297 batch_normalization_91 False
298 batch_normalization_92 False
299 conv2d_93 False
300 batch_normalization_85 False
301 activation_87 False
302 activation_88 False
303 activation_91 False
304 activation_92 False
305 batch_normalization_93 False
306 activation_85 False
307 mixed9_1 False
308 concatenate 1 False

```

```
inception_model.summary()
```

			activation
conv2d_89 (Conv2D)	(None, 5, 5, 448)	917504	mixed9[0]
batch_normalization_89 (Batch Normalization)	(None, 5, 5, 448)	1344	conv2d_89
activation_89 (Activation)	(None, 5, 5, 448)	0	batch_norm
conv2d_86 (Conv2D)	(None, 5, 5, 384)	786432	mixed9[0]
conv2d_90 (Conv2D)	(None, 5, 5, 384)	1548288	activation
batch_normalization_86 (Batch Normalization)	(None, 5, 5, 384)	1152	conv2d_86
batch_normalization_90 (Batch Normalization)	(None, 5, 5, 384)	1152	conv2d_90
activation_86 (Activation)	(None, 5, 5, 384)	0	batch_norm
activation_90 (Activation)	(None, 5, 5, 384)	0	batch_norm
conv2d_87 (Conv2D)	(None, 5, 5, 384)	442368	activation
conv2d_88 (Conv2D)	(None, 5, 5, 384)	442368	activation
conv2d_91 (Conv2D)	(None, 5, 5, 384)	442368	activation
conv2d_92 (Conv2D)	(None, 5, 5, 384)	442368	activation
average_pooling2d_8 (Average Pooling)	(None, 5, 5, 2048)	0	mixed9[0]
conv2d_85 (Conv2D)	(None, 5, 5, 320)	655360	mixed9[0]

batch_normalization_87 (BatchNo	(None, 5, 5, 384)	1152	conv2d_87[
batch_normalization_88 (BatchNo	(None, 5, 5, 384)	1152	conv2d_88[
batch_normalization_91 (BatchNo	(None, 5, 5, 384)	1152	conv2d_91[
batch_normalization_92 (BatchNo	(None, 5, 5, 384)	1152	conv2d_92[
conv2d_93 (Conv2D)	(None, 5, 5, 192)	393216	average_po
batch_normalization_85 (BatchNo	(None, 5, 5, 320)	960	conv2d_85[
activation_87 (Activation)	(None, 5, 5, 384)	0	batch_norm
activation_88 (Activation)	(None, 5, 5, 384)	0	batch_norm
activation_91 (Activation)	(None, 5, 5, 384)	0	batch_norm
activation_92 (Activation)	(None, 5, 5, 384)	0	batch_norm
batch_normalization_93 (BatchNo	(None, 5, 5, 192)	576	conv2d_93[
activation_85 (Activation)	(None, 5, 5, 320)	0	batch_norm
mixed9_1 (Concatenate)	(None, 5, 5, 768)	0	activation
			activation

#adding layers after non-trainable layers of InceptionV3

```
x = inception_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(512, activation='relu', kernel_initializer='HeUniform')(x)
x = Dense(64, activation='relu', kernel_initializer='HeUniform')(x)
x = (Dense(1, activation='sigmoid'))(x)
```

```
model3 = Model(inputs=inception_model.input, outputs=x)
```

#tensorboard callback

```
tbCallBack3 = tf.keras.callbacks.TensorBoard(log_dir='./log3', histogram_freq=0,
                                              write_graph=True,
                                              write_grads=True,
                                              write_images=True)
callback_list3 = [tbCallBack3]
```

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `1

#compiling the model

```
model3.compile(optimizer='adam', loss="binary_crossentropy", metrics=["accuracy"])
```

#training the model, only the layers below non-trainable layers will be trained

```
model3.fit_generator(train_ImageGenerator, epochs = 20, validation_data=(test_Imag
```

Epoch 1/20

```
125/125 [=====] - 159s 1s/step - loss: 0.7430 - accu
Epoch 2/20
125/125 [=====] - 142s 1s/step - loss: 0.6945 - accu
Epoch 3/20
125/125 [=====] - 142s 1s/step - loss: 0.6928 - accu
Epoch 4/20
125/125 [=====] - 127s 947ms/step - loss: 0.6830 - a
Epoch 5/20
125/125 [=====] - 129s 964ms/step - loss: 0.6835 - a
Epoch 6/20
125/125 [=====] - 144s 1s/step - loss: 0.6776 - accu
Epoch 7/20
125/125 [=====] - 127s 955ms/step - loss: 0.6729 - a
Epoch 8/20
125/125 [=====] - 128s 968ms/step - loss: 0.6723 - a
Epoch 9/20
125/125 [=====] - 142s 1s/step - loss: 0.6648 - accu
Epoch 10/20
125/125 [=====] - 128s 958ms/step - loss: 0.6730 - a
Epoch 11/20
125/125 [=====] - 127s 947ms/step - loss: 0.6637 - a
Epoch 12/20
125/125 [=====] - 127s 953ms/step - loss: 0.6614 - a
Epoch 13/20
125/125 [=====] - 127s 955ms/step - loss: 0.6659 - a
Epoch 14/20
125/125 [=====] - 127s 962ms/step - loss: 0.6666 - a
Epoch 15/20
125/125 [=====] - 128s 954ms/step - loss: 0.6577 - a
Epoch 16/20
125/125 [=====] - 128s 965ms/step - loss: 0.6514 - a
Epoch 17/20
125/125 [=====] - 129s 972ms/step - loss: 0.6572 - a
Epoch 18/20
125/125 [=====] - 133s 959ms/step - loss: 0.6537 - a
Epoch 19/20
125/125 [=====] - 127s 960ms/step - loss: 0.6591 - a
Epoch 20/20
125/125 [=====] - 141s 1s/step - loss: 0.6584 - accu
<tensorflow.python.keras.callbacks.History at 0x7fe2466fc2d0>
```



```
#loading tensorboard
%tensorboard --logdir log3
```

TensorBoard

SCALARS

GRAPHS

INACTIVE

- ☐ Show data download links
- ☐ Ignore outliers in chart scaling

Tooltip sorting method: **default** ▼

Smoothing



0.6

Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

☐ ☐ train

☐ ☐ dev

```
#predicting values for test dataset
```

```
y_pred_i = model3.predict(test_ImageGenerator)
```

```
log3
```

```
y_pred_i_df = pd.DataFrame(y_pred_i).round().astype('int')
```

```
#plotting confusion matrix
```

```
cm_i = confusion_matrix(val_df['dank_or_not'], y_pred_i_df[0])
```

```
sns.heatmap(cm_i, annot=True, fmt='g')
```

```
plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
```


```
plt.show()
```

🔍 Filter tags (regular expressions supported)

epoch_accuracy



epoch_accuracy
tag: epoch_accuracy



epoch_loss





```
#getting true negative, false positive, false negative and true positive values fr
tn_i, fp_i, fn_i, tp_i = cm_i.ravel()
```

```
    5
```

```
#computing sensetivity and specificity
sensitivity_i = (tp_i/(tp_i+fn_i)).round(4)
specificity_i = (tn_i/(tn_i+fp_i)).round(4)
print('sensitivity      : ',sensitivity_i)
print('specificity      : ',specificity_i)
print('accuracy         : ',accuracy(val_df['dank_or_not'], y_pred_i_df[0]))
```

```
sensitivity      :  0.3956
specificity      :  0.6088
accuracy         :  tf.Tensor(0.5070785, shape=(), dtype=float32)
```

```
#saving the model
model3.save('inception_model.h5')
```


▼ Applying Efficient Net model

```
#importing EfficientNetB5 model
from tensorflow.keras.applications import EfficientNetB5
```

```
%load_ext tensorboard
# Clear any logs from previous runs
!rm -rf ./log4/
```

```
#defining EfficientNetB5 model without including top fully connected layers, to pe
efficient_model = EfficientNetB5(weights='imagenet', include_top=False, input_shap
```

```
Downloading data from https://storage.googleapis.com/keras-applications/efficientnetb5/115269632/115263384 [=====] - 1s 0us/step
```



```
#setting layers of EfficientNetB5 as non-trainable
for layer in efficient_model.layers:
    layer.trainable = False
```

```
#checking layers of EfficientNetB5
for i, layer in enumerate(efficient_model.layers):
    print(i, layer.name, layer.trainable)
```

```
-----
512 block6h_project_bn False
513 block6h_drop False
514 block6h_add False
515 block6i_expand_conv False
516 block6i_expand_bn False
```



```

517 block6i_expand_activation False
518 block6i_dwconv False
519 block6i_bn False
520 block6i_activation False
521 block6i_se_squeeze False
522 block6i_se_reshape False
523 block6i_se_reduce False
524 block6i_se_expand False
525 block6i_se_excite False
526 block6i_project_conv False
527 block6i_project_bn False
528 block6i_drop False
529 block6i_add False
530 block7a_expand_conv False
531 block7a_expand_bn False
532 block7a_expand_activation False
533 block7a_dwconv False
534 block7a_bn False
535 block7a_activation False
536 block7a_se_squeeze False
537 block7a_se_reshape False
538 block7a_se_reduce False
539 block7a_se_expand False
540 block7a_se_excite False
541 block7a_project_conv False
542 block7a_project_bn False
543 block7b_expand_conv False
544 block7b_expand_bn False
545 block7b_expand_activation False
546 block7b_dwconv False
547 block7b_bn False
548 block7b_activation False
549 block7b_se_squeeze False
550 block7b_se_reshape False
551 block7b_se_reduce False
552 block7b_se_expand False

553 block7b_se_excite False
554 block7b_project_conv False
555 block7b_project_bn False
556 block7b_drop False
557 block7b_add False
558 block7c_expand_conv False
559 block7c_expand_bn False
560 block7c_expand_activation False
561 block7c_dwconv False
562 block7c_bn False
563 block7c_activation False
564 block7c_se_squeeze False
565 block7c_se_reshape False
566 block7c_se_reduce False
567 block7c_se_expand False
568 block7c_se_excite False
569 block7c_project_conv False
570 block7c_project_bn False

```

```
efficient_model.summary()
```

```

block7b_dwconv (DepthwiseConv2D (None, 7, 7, 3072)) 27648 block7b_ex

```

block7b_bn (BatchNormalization)	(None, 7, 7, 3072)	12288	block7b_dw
block7b_activation (Activation)	(None, 7, 7, 3072)	0	block7b_bn
block7b_se_squeeze (GlobalAveragePooling2D)	(None, 3072)	0	block7b_ac
block7b_se_reshape (Reshape)	(None, 1, 1, 3072)	0	block7b_se
block7b_se_reduce (Conv2D)	(None, 1, 1, 128)	393344	block7b_se
block7b_se_expand (Conv2D)	(None, 1, 1, 3072)	396288	block7b_se
block7b_se_excite (Multiply)	(None, 7, 7, 3072)	0	block7b_ac block7b_se
block7b_project_conv (Conv2D)	(None, 7, 7, 512)	1572864	block7b_se
block7b_project_bn (BatchNormalization)	(None, 7, 7, 512)	2048	block7b_pr
block7b_drop (Dropout)	(None, 7, 7, 512)	0	block7b_pr
block7b_add (Add)	(None, 7, 7, 512)	0	block7b_dr block7a_pr
block7c_expand_conv (Conv2D)	(None, 7, 7, 3072)	1572864	block7b_ad
block7c_expand_bn (BatchNormalization)	(None, 7, 7, 3072)	12288	block7c_ex
block7c_expand_activation (Activation)	(None, 7, 7, 3072)	0	block7c_ex
block7c_dwconv (DepthwiseConv2D)	(None, 7, 7, 3072)	27648	block7c_ex
block7c_bn (BatchNormalization)	(None, 7, 7, 3072)	12288	block7c_dw
block7c_activation (Activation)	(None, 7, 7, 3072)	0	block7c_bn
block7c_se_squeeze (GlobalAveragePooling2D)	(None, 3072)	0	block7c_ac
block7c_se_reshape (Reshape)	(None, 1, 1, 3072)	0	block7c_se
block7c_se_reduce (Conv2D)	(None, 1, 1, 128)	393344	block7c_se
block7c_se_expand (Conv2D)	(None, 1, 1, 3072)	396288	block7c_se
block7c_se_excite (Multiply)	(None, 7, 7, 3072)	0	block7c_ac block7c_se
block7c_project_conv (Conv2D)	(None, 7, 7, 512)	1572864	block7c_se
block7c_project_bn (BatchNormalization)	(None, 7, 7, 512)	2048	block7c_pr
block7c_drop (Dropout)	(None, 7, 7, 512)	0	block7c_pr
block7c_add (Add)	(None, 7, 7, 512)	0	block7c_dr block7b_ad

#adding layers after non-trainable layers of EfficientNetB5

```
x = efficient_model.output
```

```
x = GlobalAveragePooling2D()(x)
```

```
x = Dense(512, activation='relu', kernel_initializer='HeUniform')(x)
```

```


x = (Dense(1, activation='sigmoid'))(x)

model4 = Model(inputs=efficient_model.input, outputs=x)

#tensorboard callback
tbCallBack4 = tf.keras.callbacks.TensorBoard(log_dir='./log4', histogram_freq=0,
                                             write_graph=True,
                                             write_grads=True,
                                             write_images=True)
callback_list4 = [tbCallBack4]

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `1

```



```

#compiling the model
model4.compile(optimizer='adam', loss="binary_crossentropy", metrics=["accuracy"])

#training the model, only the layers below non-trainable layers will be trained
model4.fit_generator(train_ImageGenerator, epochs = 20, validation_data=(test_Imag

```

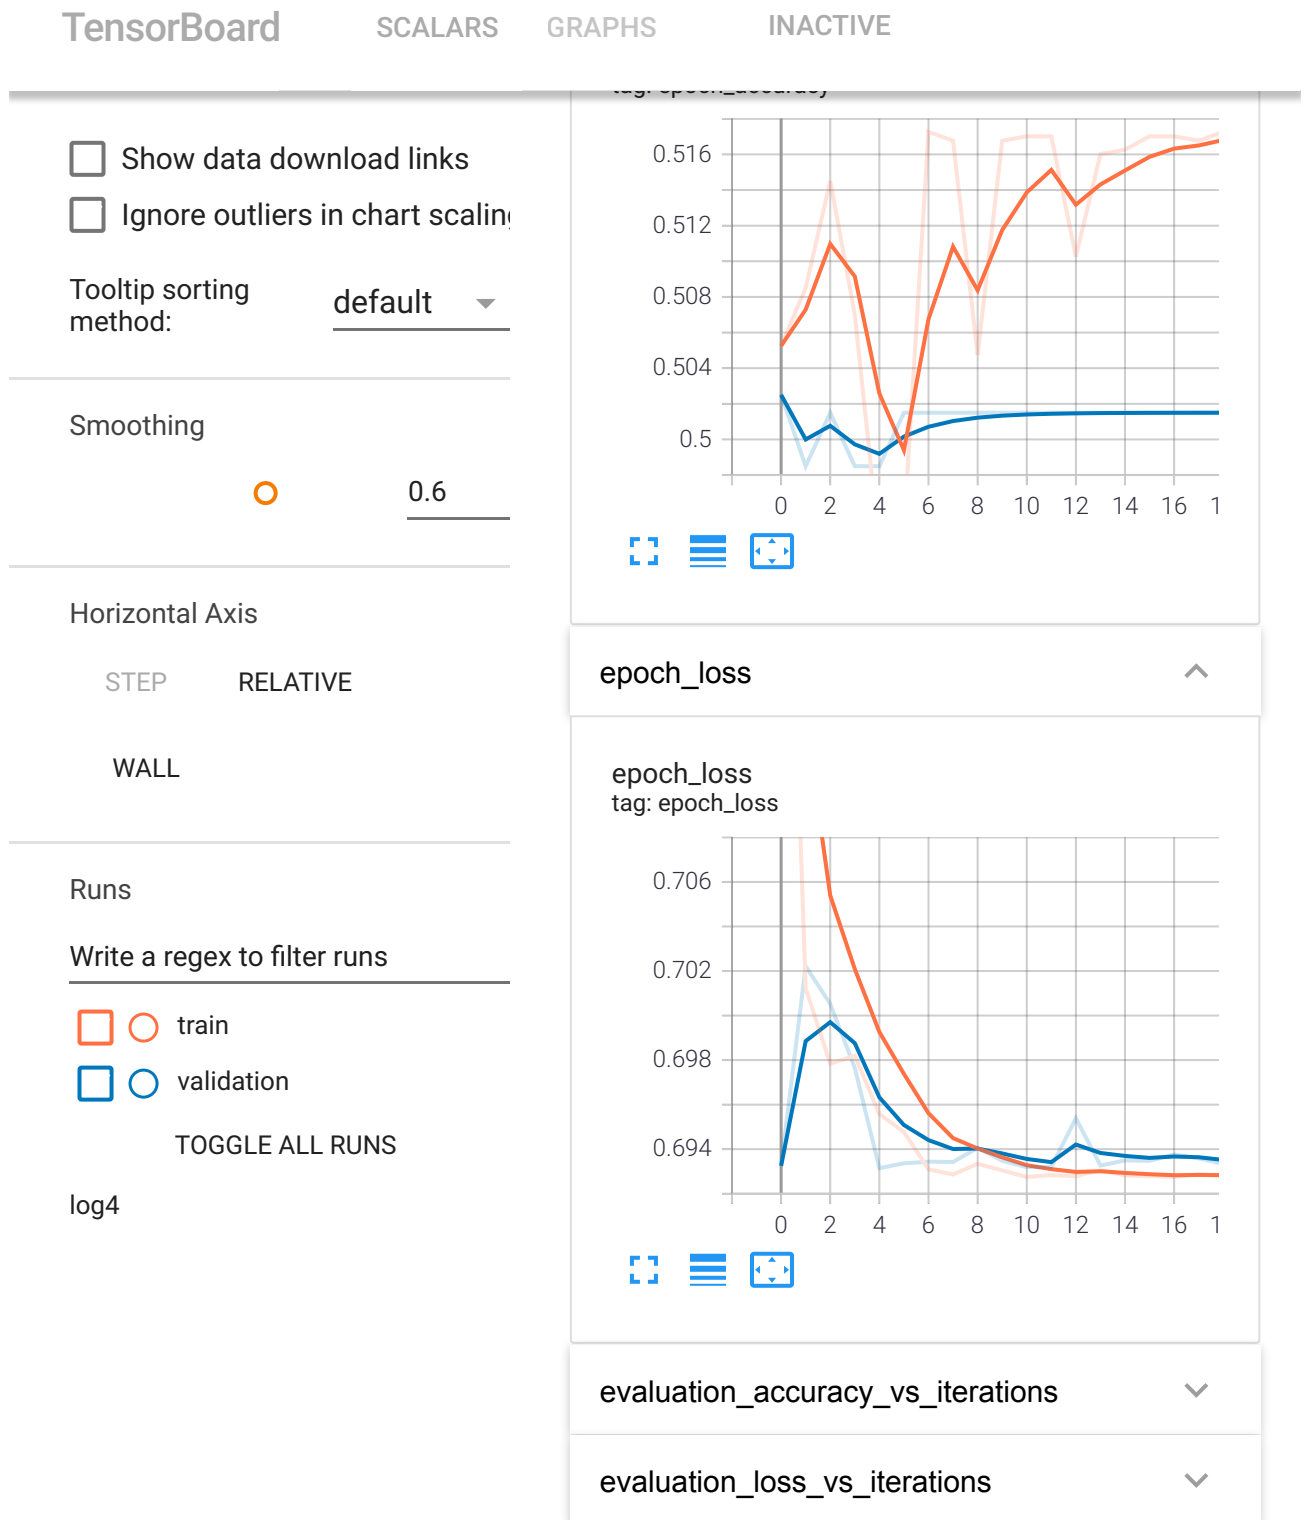
```

Epoch 1/20
125/125 [=====] - 233s 1s/step - loss: 0.7333 - accu
Epoch 2/20
125/125 [=====] - 132s 996ms/step - loss: 0.7012 - a
Epoch 3/20
125/125 [=====] - 146s 1s/step - loss: 0.6978 - accu
Epoch 4/20
125/125 [=====] - 145s 1s/step - loss: 0.6982 - accu
Epoch 5/20
125/125 [=====] - 132s 999ms/step - loss: 0.6956 - a
Epoch 6/20
125/125 [=====] - 145s 1s/step - loss: 0.6948 - accu
Epoch 7/20
125/125 [=====] - 131s 983ms/step - loss: 0.6931 - a
Epoch 8/20
125/125 [=====] - 129s 966ms/step - loss: 0.6929 - a
Epoch 9/20
125/125 [=====] - 130s 968ms/step - loss: 0.6933 - a
Epoch 10/20
125/125 [=====] - 130s 983ms/step - loss: 0.6930 - a
Epoch 11/20
125/125 [=====] - 144s 1s/step - loss: 0.6928 - accu
Epoch 12/20
125/125 [=====] - 131s 976ms/step - loss: 0.6928 - a
Epoch 13/20
125/125 [=====] - 135s 985ms/step - loss: 0.6928 - a
Epoch 14/20
125/125 [=====] - 143s 1s/step - loss: 0.6931 - accu
Epoch 15/20
125/125 [=====] - 131s 981ms/step - loss: 0.6928 - a
Epoch 16/20
125/125 [=====] - 145s 1s/step - loss: 0.6928 - accu
Epoch 17/20
125/125 [=====] - 145s 1s/step - loss: 0.6928 - accu
Epoch 18/20
125/125 [=====] - 129s 964ms/step - loss: 0.6929 - a
Epoch 19/20

```

```
125/125 [=====] - 150s 1s/step - loss: 0.6928 - accu
Epoch 20/20
125/125 [=====] - 130s 981ms/step - loss: 0.6926 - a
<tensorflow.python.keras.callbacks.History at 0x7fe25dc02210>
```

```
#loading the tensorboard
%tensorboard --logdir log4
```




```
#predicting the test data
y_pred_e = model4.predict(test_ImageGenerator)
```

```

y_pred_e_df = pd.DataFrame(y_pred_e).round().astype('int')

#plotting confusion matrix
cm_e = confusion_matrix(val_df['dank_or_not'], y_pred_e_df[0])
sns.heatmap(cm_e, annot=True, fmt='g')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```



```

#getting true negative, false positive, false negative and true positive values fr
tn_e, fp_e, fn_e, tp_e = cm_e.ravel()

```

```

#computing sensetivity and specificity
sensitivity_e = (tp_e/(tp_e+fn_e)).round(4)
specificity_e = (tn_e/(tn_e+fp_e)).round(4)
print('sensitivity      : ',sensitivity_e)
print('specificity      : ',specificity_e)
print('accuracy         : ',accuracy(val_df['dank_or_not'], y_pred_e_df[0]))

```

```

sensitivity      :  0.0
specificity      :  1.0
accuracy         :  tf.Tensor(0.5063814, shape=(), dtype=float32)

```

```

#saving the model
model4.save('efficientnet_model.h5')

```

▼ Applying Resnet50

```

#importing ResNet50 model
from tensorflow.keras.applications.resnet50 import ResNet50

```

```

%load_ext tensorboard
# Clear any logs from previous runs


```

```
!rm -rf ./log5/
```

```
The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard
```

```
#defining ResNet50 model without including top fully connected layers, to perform
resnet_model = ResNet50(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50\_weights\_tf\_dim\_ordering\_tf\_kernels\_notop.h5
94773248/94765736 [=====] - 1s 0us/step
```



```
#setting layers of ResNet50 as non-trainable
for layer in resnet_model.layers:
    layer.trainable = False
```

```
#checking layers of ResNet50
for i, layer in enumerate(resnet_model.layers):
    print(i, layer.name, layer.trainable)
```

```
115 conv4_block4_1_relu False
116 conv4_block4_2_conv False
117 conv4_block4_2_bn False
118 conv4_block4_2_relu False
119 conv4_block4_3_conv False
120 conv4_block4_3_bn False
121 conv4_block4_add False
122 conv4_block4_out False
123 conv4_block5_1_conv False
124 conv4_block5_1_bn False
125 conv4_block5_1_relu False
126 conv4_block5_2_conv False
127 conv4_block5_2_bn False
128 conv4_block5_2_relu False
129 conv4_block5_3_conv False
130 conv4_block5_3_bn False
131 conv4_block5_add False
132 conv4_block5_out False
133 conv4_block6_1_conv False
134 conv4_block6_1_bn False
135 conv4_block6_1_relu False
136 conv4_block6_2_conv False
137 conv4_block6_2_bn False
138 conv4_block6_2_relu False
139 conv4_block6_3_conv False
140 conv4_block6_3_bn False
141 conv4_block6_add False
142 conv4_block6_out False
143 conv5_block1_1_conv False
144 conv5_block1_1_bn False
145 conv5_block1_1_relu False
146 conv5_block1_2_conv False
147 conv5_block1_2_bn False
148 conv5_block1_2_relu False
149 conv5_block1_0_conv False
150 conv5_block1_3_conv False
151 conv5_block1_0_bn False
152 conv5_block1_3_bn False
```

```

152 conv5_block1_2_relu (Activation) (None, 7, 7, 2048) 0 conv5_block1_2_relu
153 conv5_block1_add (Add) (None, 7, 7, 2048) 0 conv5_block1_add
154 conv5_block1_out (Activation) (None, 7, 7, 2048) 0 conv5_block1_out
155 conv5_block2_1_conv (Conv2D) (None, 7, 7, 512) 1049088 conv5_block2_1_conv
156 conv5_block2_1_bn (BatchNormali (None, 7, 7, 512) 2048 conv5_block2_1_bn
157 conv5_block2_1_relu (Activation) (None, 7, 7, 512) 0 conv5_block2_1_relu
158 conv5_block2_2_conv (Conv2D) (None, 7, 7, 512) 2359808 conv5_block2_2_conv
159 conv5_block2_2_bn (BatchNormali (None, 7, 7, 512) 2048 conv5_block2_2_bn
160 conv5_block2_2_relu (Activation) (None, 7, 7, 512) 0 conv5_block2_2_relu
161 conv5_block2_3_conv (Conv2D) (None, 7, 7, 2048) 1050624 conv5_block2_3_conv
162 conv5_block2_3_bn (BatchNormali (None, 7, 7, 2048) 8192 conv5_block2_3_bn
163 conv5_block2_add (Add) (None, 7, 7, 2048) 0 conv5_block2_add
164 conv5_block2_out (Activation) (None, 7, 7, 2048) 0 conv5_block2_out
165 conv5_block3_1_conv (Conv2D) (None, 7, 7, 2048) 1050624 conv5_block3_1_conv
166 conv5_block3_1_bn (BatchNormali (None, 7, 7, 2048) 8192 conv5_block3_1_bn
167 conv5_block3_1_relu (Activation) (None, 7, 7, 2048) 0 conv5_block3_1_relu
168 conv5_block3_2_conv (Conv2D) (None, 7, 7, 2048) 1050624 conv5_block3_2_conv
169 conv5_block3_2_bn (BatchNormali (None, 7, 7, 2048) 8192 conv5_block3_2_bn
170 conv5_block3_2_relu (Activation) (None, 7, 7, 2048) 0 conv5_block3_2_relu
171 conv5_block3_3_conv (Conv2D) (None, 7, 7, 2048) 1050624 conv5_block3_3_conv
172 conv5_block3_3_bn (BatchNormali (None, 7, 7, 2048) 8192 conv5_block3_3_bn
173 conv5_block3_add (Add) (None, 7, 7, 2048) 0 conv5_block3_add

```

```
resnet_model.summary()
```

conv5_block1_2_relu (Activation)	(None, 7, 7, 512)	0	conv5_block1_2_relu
conv5_block1_0_conv (Conv2D)	(None, 7, 7, 2048)	2099200	conv4_block1_out
conv5_block1_3_conv (Conv2D)	(None, 7, 7, 2048)	1050624	conv5_block1_3_conv
conv5_block1_0_bn (BatchNormali	(None, 7, 7, 2048)	8192	conv5_block1_0_bn
conv5_block1_3_bn (BatchNormali	(None, 7, 7, 2048)	8192	conv5_block1_3_bn
conv5_block1_add (Add)	(None, 7, 7, 2048)	0	conv5_block1_add conv5_block1_out
conv5_block1_out (Activation)	(None, 7, 7, 2048)	0	conv5_block1_out
conv5_block2_1_conv (Conv2D)	(None, 7, 7, 512)	1049088	conv5_block2_1_conv
conv5_block2_1_bn (BatchNormali	(None, 7, 7, 512)	2048	conv5_block2_1_bn
conv5_block2_1_relu (Activation)	(None, 7, 7, 512)	0	conv5_block2_1_relu
conv5_block2_2_conv (Conv2D)	(None, 7, 7, 512)	2359808	conv5_block2_2_conv
conv5_block2_2_bn (BatchNormali	(None, 7, 7, 512)	2048	conv5_block2_2_bn
conv5_block2_2_relu (Activation)	(None, 7, 7, 512)	0	conv5_block2_2_relu
conv5_block2_3_conv (Conv2D)	(None, 7, 7, 2048)	1050624	conv5_block2_3_conv
conv5_block2_3_bn (BatchNormali	(None, 7, 7, 2048)	8192	conv5_block2_3_bn
conv5_block2_add (Add)	(None, 7, 7, 2048)	0	conv5_block2_add conv5_block2_out
conv5_block2_out (Activation)	(None, 7, 7, 2048)	0	conv5_block2_out

conv5_block3_1_conv (Conv2D)	(None, 7, 7, 512)	1049088	conv5_block3_1_conv
conv5_block3_1_bn (BatchNormali	(None, 7, 7, 512)	2048	conv5_block3_1_bn
conv5_block3_1_relu (Activation	(None, 7, 7, 512)	0	conv5_block3_1_relu
conv5_block3_2_conv (Conv2D)	(None, 7, 7, 512)	2359808	conv5_block3_2_conv
conv5_block3_2_bn (BatchNormali	(None, 7, 7, 512)	2048	conv5_block3_2_bn
conv5_block3_2_relu (Activation	(None, 7, 7, 512)	0	conv5_block3_2_relu
conv5_block3_3_conv (Conv2D)	(None, 7, 7, 2048)	1050624	conv5_block3_3_conv
conv5_block3_3_bn (BatchNormali	(None, 7, 7, 2048)	8192	conv5_block3_3_bn
conv5_block3_add (Add)	(None, 7, 7, 2048)	0	conv5_block3_add
conv5_block3_out (Activation)	(None, 7, 7, 2048)	0	conv5_block3_out
=====			
Total params: 23,587,712			

```
#adding layers after non-trainable layers of ResNet50
```

```
x = resnet_model.output
```

```
x = GlobalAveragePooling2D()(x)
```

```
x = Dense(512, activation='relu', kernel_initializer='HeUniform')(x)
```

```
x = (Dense(1, activation='sigmoid'))(x)
```

```
model5 = Model(inputs=resnet_model.input, outputs=x)
```

```
#tensorboard callback
```

```
tbCallBack5 = tf.keras.callbacks.TensorBoard(log_dir='./log5', histogram_freq=0,
                                              write_graph=True,
                                              write_grads=True,
                                              write_images=True)
```

```
callback_list5 = [tbCallBack5]
```

```
WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `1
```

```
#compiling the model
```

```
model5.compile(optimizer='adam', loss="binary_crossentropy", metrics=["accuracy"])
```

```
#training the model, only the layers below non-trainable layers will be trained
```

```
model5.fit_generator(train_ImageGenerator, epochs = 20, validation_data=(test_Imag
```

```
Epoch 1/20
```

```
125/125 [=====] - 135s 992ms/step - loss: 0.7400 - a
```

```
Epoch 2/20
```

```
125/125 [=====] - 128s 973ms/step - loss: 0.7046 - a
```

```
Epoch 3/20
```


```
125/125 [=====] - 128s 976ms/step - loss: 0.7197 - a
```

```
Epoch 4/20
```

```

125/125 [=====] - 129s 962ms/step - loss: 0.7088 - a
Epoch 5/20
125/125 [=====] - 128s 969ms/step - loss: 0.6980 - a
Epoch 6/20
125/125 [=====] - 129s 963ms/step - loss: 0.6938 - a
Epoch 7/20
125/125 [=====] - 129s 960ms/step - loss: 0.6921 - a
Epoch 8/20
125/125 [=====] - 129s 978ms/step - loss: 0.6924 - a
Epoch 9/20
125/125 [=====] - 146s 1s/step - loss: 0.6897 - accu
Epoch 10/20
125/125 [=====] - 132s 984ms/step - loss: 0.6943 - a
Epoch 11/20
125/125 [=====] - 135s 969ms/step - loss: 0.6884 - a
Epoch 12/20
125/125 [=====] - 128s 966ms/step - loss: 0.6883 - a
Epoch 13/20
125/125 [=====] - 128s 953ms/step - loss: 0.6881 - a
Epoch 14/20
125/125 [=====] - 129s 970ms/step - loss: 0.6859 - a
Epoch 15/20
125/125 [=====] - 129s 967ms/step - loss: 0.6894 - a
Epoch 16/20
125/125 [=====] - 130s 973ms/step - loss: 0.6872 - a
Epoch 17/20
125/125 [=====] - 131s 972ms/step - loss: 0.6883 - a
Epoch 18/20
125/125 [=====] - 145s 1s/step - loss: 0.6891 - accu
Epoch 19/20
125/125 [=====] - 132s 989ms/step - loss: 0.6889 - a
Epoch 20/20
125/125 [=====] - 133s 1s/step - loss: 0.6886 - accu
<tensorflow.python.keras.callbacks.History at 0x7fe30cec8090>

```



```

#loading the tensorboard
%tensorboard --logdir log5

```

TensorBoard

SCALARS

GRAPHS

INACTIVE

☐ Show data download links☐ Ignore outliers in chart scalingTooltip sorting
method: **default**

Smoothing



0.6

Horizontal Axis


STEP

RELATIVE

WALL

Filter tags (regular expressions supported)

epoch_accuracy

epoch_accuracy
tag: epoch_accuracy

#predicting the test data

y_pred_r = model5.predict(test_ImageGenerator)

y_pred_r_df = pd.DataFrame(y_pred_r).round().astype('int')

#plotting confusion matrix


cm_r = confusion_matrix(val_df['dank_or_not'], y_pred_r_df[0])

sns.heatmap(cm_r, annot=True, fmt='g')

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.show()



#getting true negative, false positive, false negative and true positive values fr

```

tn_r, tp_r, tn_r, fp_r = cm_r.ravel()

#computing sensitivity and specificity
sensitivity_r = (tp_r/(tp_r+fn_r)).round(4)
specificity_r = (tn_r/(tn_r+fp_r)).round(4)
print('sensitivity      : ',sensitivity_r)
print('specificity      : ',specificity_r)
print('accuracy         : ',accuracy(val_df['dank_or_not'], y_pred_r_df[0]))

sensitivity      :  0.4819
specificity      :  0.489
accuracy         :  tf.Tensor(0.50158495, shape=(), dtype=float32)

#saving the model
model5.save('resnet_model.h5')

```

▼ Comparing all models

```

from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Model Name", "Sensitivity", "Specificity"]
x.add_row(["Baseline", 0.50, 0.51])
x.add_row(["VGG16", 0.4719, 0.5269])
x.add_row(["InceptionV3", 0.3956, 0.6088])
x.add_row(["EfficientNetB5", 0.0, 1.0])
x.add_row(["ResNet50", 0.4819, 0.489])

print(x)

```

Model Name	Sensitivity	Specificity
Baseline	0.5	0.51
VGG16	0.4719	0.5269
InceptionV3	0.3956	0.6088
EfficientNetB5	0.0	1.0
ResNet50	0.4819	0.489

From the table we can see that ResNet50 and VGG16 are best, InceptionV3 is also performing well and EfficientNetB5 is worst.

✓ 0s completed at 12:32 PM

● ✕